

BFG Repo-Cleaner

Removes large or troublesome blobs like `git-filter-branch` does, but faster. And written in Scala



```
$ bfg --strip-blobs-bigger-than 100M --replace-text banned.txt repo.git
```

an alternative to git-filter-branch

The BFG is a simpler, faster alternative to `git-filter-branch` for cleansing bad data out of your Git repository history:

- Removing **Crazy Big Files**
- Removing **Passwords, Credentials** & other **Private data**

The `git-filter-branch` command is enormously powerful and can do things that the BFG can't - but the BFG is *much* better for the tasks above, because:

- **Faster** : **10 - 720x** faster
- **Simpler** : The BFG isn't particularly clever, but *is* focused on making the above tasks easy
- **Beautiful** : If you need to, you can use the beautiful Scala language to customise the BFG. Which has got to be better than Bash scripting at least some of the time.

Usage

First clone a fresh copy of your repo, using the `--mirror` flag:

```
$ git clone --mirror git://example.com/some-big-repo.git
```

This is a **bare** repo, which means your normal files won't be visible, but it is a *full* copy of the Git database of your repository, and at this point you should **make a backup of it** to ensure you don't lose anything.

Now you can run the BFG to clean your repository up:

```
$ java -jar bfg.jar --strip-blobs-bigger-than 100M some-big-repo.git
```

The BFG will update your commits and all branches and tags so they are clean, but it doesn't physically delete the unwanted stuff. Examine the repo to make sure your history has been updated, and then use the standard `git gc` command to strip out the unwanted dirty data, which Git will now recognise as surplus to requirements:

```
$ cd some-big-repo.git
$ git reflog expire --expire=now --all && git gc --prune=now --aggressive
```

Finally, once you're happy with the updated state of your repo, push it back up (*note that because your clone command used the `--mirror` flag, this push will update all refs on your remote server*):

```
$ git push
```

At this point, you're ready for everyone to ditch their old copies of the repo and do fresh clones of the nice, new pristine data. It's best to delete all old clones, as they'll have dirty history that you *don't* want to risk pushing back into your newly cleaned repo.

Examples

In all these examples `bfg` is an alias for `java -jar bfg.jar`.

Delete all files named 'id_rsa' or 'id_dsa' :

```
$ bfg --delete-files id_{dsa,rsa} my-repo.git
```

Remove all blobs bigger than 50 megabytes :

```
$ bfg --strip-blobs-bigger-than 50M my-repo.git
```

Replace all passwords listed in a file (*prefix lines 'regex:' or 'glob:' if required*) with `***REMOVED***` wherever they occur in your repository :

```
$ bfg --replace-text passwords.txt my-repo.git
```

Remove all folders or files named '.git' - a **reserved filename** in Git. These often **become a problem** when migrating to Git from other source-control systems like Mercurial :

```
$ bfg --delete-folders .git --delete-files .git --no-blob-protection my-repo.git
```

For further command-line options, you can run the BFG without any arguments, which will output [text like this](#).

Your current files are sacred...

By default the BFG doesn't modify the contents of your *latest* commit on your `master` (or `HEAD`) branch, even though it *will* clean all the commits before it.

That's because your latest commit is likely to be the one that you deploy to production, and a simple deletion of a private credential or a big file is quite likely to result in broken code that no longer has the hard-coded data it expects - you need to fix that, the BFG can't do it for you. Once you've committed your changes- and your latest commit is *clean* with none of the undesired data in it - you can run the BFG to perform it's simple deletion operations over all your historical commits.

Note:

- Cleaning Git repos is about *completely* eradicating bad stuff from history. If something 'bad' (like a 10MB file, when you're specifying `--strip-blobs-bigger-than 5M`) is in a protected commit, it *won't* be deleted - it'll persist in your repository, **even if the BFG deletes if from earlier commits**. If you want the BFG to delete something **you need to make sure your current commits are clean**.
- Note that although the files in those protected commits won't be changed, when those commits follow on from earlier dirty commits, their commit ids **will** change, to reflect the changed history - only the SHA-1 id of the filesystem-tree will remain the same.

If you want to turn off the protection (in general, not recommended) you can use the `--no-blob-protection` flag:

```
$ bfg --strip-biggest-blobs 100 --no-blob-protection repo.git
```

Faster...

The BFG is **10 - 720x** faster than `git-filter-branch`, turning an *overnight* job into one that takes *less than ten minutes*.

BFG's performance advantage is due to these factors:

- The approach of `git-filter-branch` is to step through every commit in your repository, examining the complete file-hierarchy of each one. For the intended use-cases of The BFG this is wasteful, as we don't care *where* in a file structure a 'bad' file exists - we just want it dealt with. Inherent in the nature of Git is that **every** file and folder is represented precisely once (and given a unique **SHA-1** hash-id). The BFG takes advantage of this to process each and every file & folder exactly **once** - no need for extra work.
- Taking advantage of the great support for parallelism in **Scala** and the JVM, the BFG does multi-core processing by default - the work of cleaning your Git repository is spread over every single core in your machine and typically consumes 100% of capacity for a substantial portion of the run.
- All action takes place in a single process (the process of the JVM), so doesn't require the frequent fork-and-exec-ing needed by `git-filter-branch`'s mix of Bash and C code.

Feedback

I tried deleting using several "how to" blog entries for git filter-branch, but wasn't successful. Then tried The BFG; worked like a champ - very cool tool!

— [Bill Hunt](#), CTO at [OptTown](#)

I found The BFG Repo-Cleaner and ran it to clean up some large files, and was amazed by the performance.

— [Jason Frey](#), Software Engineer at [Red Hat](#)

I was able to shrink the current repository down to ~500 megabytes in about 10 minutes when using this tool. My hand crafted scripts clock in at 615 megabytes in 3 days time for comparison.

— [Elliot Glaysher](#), Google Software Engineer on [Google Chrome](#)

The BFG was simple to set up and so fast that I had to ask Roberto, *"Is that it?"* and check for myself... it worked exactly as intended.

— [Nicholas Tollervey](#), Developer at [The Guardian](#)

Roberto's creations ([Agit](#) and The BFG) are both very cool ;-)

— [Junio C Hamano](#), Maintainer of [Git](#)

Also see more feedback on [Twitter](#)...

Requirements

- The [Java Runtime Environment](#) (Java 8 or above - BFG [v1.12.16](#) was the last version to support Java 7)

That's it - the Scala library and all other dependencies are folded into the [downloadable jar](#).


Links...

- [Rewriting Git project history with The BFG](#) - a blogpost for The Guardian
- [GitMinutes](#) podcast interview
- [Git Going Faster... with Scala](#) - talk for ScalaDays 2014, later Parleys [Presentation of the Day](#)
- [InfoQ interview](#)
- [Questions tagged git-rewrite-history](#) on Stack Overflow

License

The BFG is free software: you can redistribute it and/or modify it under the terms of the [GNU General Public License](#) as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

The BFG is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.



The BFG is by Roberto Tyley, the author of [Prout](#), [gu:who](#), [Agit](#) and the packager of [Spongy Castle](#).
[Twitter](#) [Google+](#) [PGP](#)



This page was generated by [GitHub Pages](#) using the [Architect](#) theme by [Jason Long](#).